

[About Keras](#)
[Getting started](#)
[Developer guides](#)
[Keras API reference](#)
[Models API](#)
[Layers API](#)
[Callbacks API](#)
[Optimizers](#)
[Metrics](#)
[Losses](#)
[Data loading](#)
[Built-in small datasets](#)
[Keras Applications](#)
[Mixed precision](#)
[Utilities](#)
[KerasTuner](#)
[Code examples](#)
[Why choose Keras?](#)
[Community & governance](#)
[Contributing to Keras](#)
[KerasTuner](#)
[» Keras API reference](#) / [Optimizers](#) / SGD

SGD

SGD class

```
tf.keras.optimizers.SGD(
    learning_rate=0.01, momentum=0.0, nesterov=False, name="SGD", **kwargs
)
```

Gradient descent (with momentum) optimizer.

Update rule for parameter w with gradient g when momentum is 0:

$$w = w - \text{learning_rate} * g$$

Update rule when momentum is larger than 0:

```
velocity = momentum * velocity - learning_rate * g
w = w + velocity
```

When $\text{nesterov}=\text{True}$, this rule becomes:

```
velocity = momentum * velocity - learning_rate * g
w = w + momentum * velocity - learning_rate * g
```

Arguments

- **learning_rate:** A [Tensor](#), floating point value, or a schedule that is a [tf.keras.optimizers.schedules.LearningRateSchedule](#), or a callable that takes no arguments and returns the actual value to use. The learning rate. Defaults to 0.01.
- **momentum:** float hyperparameter ≥ 0 that accelerates gradient descent in the relevant direction and dampens oscillations. Defaults to 0, i.e., vanilla gradient descent.
- **nesterov:** boolean. Whether to apply Nesterov momentum. Defaults to [False](#).
- **name:** Optional name prefix for the operations created when applying gradients. Defaults to "SGD".
- ****kwargs:** Keyword arguments. Allowed to be one of "clipnorm" or "clipvalue". "clipnorm" (float) clips gradients by norm; "clipvalue" (float) clips gradients by value.

Usage:

```
>>> opt = tf.keras.optimizers.SGD(learning_rate=0.1)
>>> var = tf.Variable(1.0)
>>> loss = lambda: (var ** 2)/2.0 # d(loss)/d(var1) = var1
>>> step_count = opt.minimize(loss, [var]).numpy()
>>> # Step is `learning_rate * grad`
>>> var.numpy()
0.9
```

```
>>> opt = tf.keras.optimizers.SGD(learning_rate=0.1, momentum=0.9)
>>> var = tf.Variable(1.0)
>>> val0 = var.value()
>>> loss = lambda: (var ** 2)/2.0          # d(loss)/d(var1) = var1
>>> # First step is '- learning_rate * grad'
>>> step_count = opt.minimize(loss, [var]).numpy()
>>> val1 = var.value()
>>> (val0 - val1).numpy()
0.1
>>> # On later steps, step-size increases because of momentum
>>> step_count = opt.minimize(loss, [var]).numpy()
>>> val2 = var.value()
>>> (val1 - val2).numpy()
0.18
```

Reference

- For `nesterov=True`, See [Sutskever et al., 2013](#).
-