

sklearn.linear_model.Ridge

```
class sklearn.linear_model.Ridge(alpha=1.0, *, fit_intercept=True, normalize='deprecated', copy_X=True, max_iter=None, tol=0.001, solver='auto', positive=False, random_state=None)
```

[\[source\]](#)

Linear least squares with L2 regularization.

Minimizes the objective function:

$$\|y - Xw\|^2 + \alpha \|w\|^2$$

This model solves a regression model where the loss function is the linear least squares function and regularization is given by the L2-norm. Also known as Ridge Regression or Tikhonov regularization. This estimator has built-in support for multi-variate regression (i.e., when y is a 2d-array of shape (n_samples, n_targets)).

Read more in the [User Guide](#).

Parameters:

alpha : {float, ndarray of shape (n_targets,)}, default=1.0

Regularization strength; must be a positive float. Regularization improves the conditioning of the problem and reduces the variance of the estimates. Larger values specify stronger regularization. Alpha corresponds to $1 / (2C)$ in other linear models such as [LogisticRegression](#) or [LinearSVC](#). If an array is passed, penalties are assumed to be specific to the targets. Hence they must correspond in number.

fit_intercept : bool, default=True

Whether to fit the intercept for this model. If set to false, no intercept will be used in calculations (i.e. x and y are expected to be centered).

normalize : bool, default=False

This parameter is ignored when `fit_intercept` is set to False. If True, the regressors X will be normalized before regression by subtracting the mean and dividing by the L2-norm. If you wish to standardize, please use [StandardScaler](#) before calling `fit` on an estimator with `normalize=False`.

Deprecated since version 1.0: `normalize` was deprecated in version 1.0 and will be removed in 1.2.

copy_X : bool, default=True

If True, X will be copied; else, it may be overwritten.

max_iter : int, default=None

Maximum number of iterations for conjugate gradient solver. For 'sparse_cg' and 'lsqr' solvers, the default value is determined by `scipy.sparse.linalg`. For 'sag' solver, the default value is 1000. For 'lbfgs' solver, the default value is 15000.

tol : float, default=1e-3

Precision of the solution.

solver : {'auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg', 'sag', 'saga', 'lbfgs'}, default='auto'

Solver to use in the computational routines:

- 'auto' chooses the solver automatically based on the type of data.
- 'svd' uses a Singular Value Decomposition of X to compute the Ridge coefficients. More stable for singular matrices than 'cholesky'.
- 'cholesky' uses the standard `scipy.linalg.solve` function to obtain a closed-form solution.

- ‘sparse_cg’ uses the conjugate gradient solver as found in `scipy.sparse.linalg.cg`. As an iterative algorithm, this solver is more appropriate than ‘cholesky’ for large-scale data (possibility to set `tol` and `max_iter`).
- ‘lsqr’ uses the dedicated regularized least-squares routine `scipy.sparse.linalg.lsqr`. It is the fastest and uses an iterative procedure.
- ‘sag’ uses a Stochastic Average Gradient descent, and ‘saga’ uses its improved, unbiased version named SAGA. Both methods also use an iterative procedure, and are often faster than other solvers when both `n_samples` and `n_features` are large. Note that ‘sag’ and ‘saga’ fast convergence is only guaranteed on features with approximately the same scale. You can preprocess the data with a scaler from `sklearn.preprocessing`.
- ‘lbfgs’ uses L-BFGS-B algorithm implemented in `scipy.optimize.minimize`. It can be used only when `positive` is True.

All last six solvers support both dense and sparse data. However, only ‘sag’, ‘sparse_cg’, and ‘lbfgs’ support sparse input when `fit_intercept` is True.

New in version 0.17: Stochastic Average Gradient descent solver.

New in version 0.19: SAGA solver.

positive : bool, default=False

When set to `True`, forces the coefficients to be positive. Only ‘lbfgs’ solver is supported in this case.

random_state : int, RandomState instance, default=None

Used when `solver == ‘sag’` or ‘saga’ to shuffle the data. See [Glossary](#) for details.

New in version 0.17: `random_state` to support Stochastic Average Gradient.

Attributes:

coef_ : ndarray of shape (n_features,) or (n_targets, n_features)

Weight vector(s).

intercept_ : float or ndarray of shape (n_targets,)

Independent term in decision function. Set to 0.0 if `fit_intercept = False`.

n_iter_ : None or ndarray of shape (n_targets,)

Actual number of iterations for each target. Available only for sag and lsqr solvers. Other solvers will return None.

New in version 0.17.

n_features_in_ : int

Number of features seen during [fit](#).

New in version 0.24.

feature_names_in_ : ndarray of shape (n_features_in_,)

Names of features seen during [fit](#). Defined only when `x` has feature names that are all strings.

New in version 1.0.

See also:

[RidgeClassifier](#)

Ridge classifier.

[RidgeCV](#)

Ridge regression with built-in cross validation.

[KernelRidge](#)

Kernel ridge regression combines ridge regression with the kernel trick.

Examples

Toggle Menu

```
>>> from sklearn.linear_model import Ridge
>>> import numpy as np
>>> n_samples, n_features = 10, 5
>>> rng = np.random.RandomState(0)
>>> y = rng.randn(n_samples)
>>> X = rng.randn(n_samples, n_features)
>>> clf = Ridge(alpha=1.0)
>>> clf.fit(X, y)
Ridge()
```

Methods

fit (X, y[, sample_weight])	Fit Ridge regression model.
get_params ([deep])	Get parameters for this estimator.
predict (X)	Predict using the linear model.
score (X, y[, sample_weight])	Return the coefficient of determination of the prediction.
set_params (**params)	Set the parameters of this estimator.

fit(X, y, sample_weight=None)

[source]

Fit Ridge regression model.

Parameters:

X : {ndarray, sparse matrix} of shape (n_samples, n_features)

Training data.

y : ndarray of shape (n_samples,) or (n_samples, n_targets)

Target values.

sample_weight : float or ndarray of shape (n_samples,), default=None

Individual weights for each sample. If given a float, every sample will have the same weight.

Returns:

self : object

Fitted estimator.

get_params(deep=True)

[source]

Get parameters for this estimator.

Parameters:

deep : bool, default=True

If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

params : dict

Parameter names mapped to their values.

predict(X)

[source]

Predict using the linear model.

X : array-like or sparse matrix, shape (n_samples, n_features)

Samples.

Returns:

C : array, shape (n_samples,)

Returns predicted values.

```
score(X, y, sample_weight=None)
```

[\[source\]](#)

Return the coefficient of determination of the prediction.

The coefficient of determination R^2 is defined as $(1 - \frac{u}{v})$, where u is the residual sum of squares $((y_{\text{true}} - y_{\text{pred}})^2).sum()$ and v is the total sum of squares $((y_{\text{true}} - y_{\text{true}.mean()})^2).sum()$. The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of y , disregarding the input features, would get a R^2 score of 0.0.

Parameters:

X : array-like of shape (n_samples, n_features)

Test samples. For some estimators this may be a precomputed kernel matrix or a list of generic objects instead with shape $(n_{\text{samples}}, n_{\text{samples_fitted}})$, where $n_{\text{samples_fitted}}$ is the number of samples used in the fitting for the estimator.

y : array-like of shape (n_samples,) or (n_samples, n_outputs)

True values for X .

sample_weight : array-like of shape (n_samples,), default=None

Sample weights.

Returns:

score : float

R^2 of `self.predict(X)` wrt. y .

Notes

The R^2 score used when calling `score` on a regressor uses `multioutput='uniform_average'` from version 0.23 to keep consistent with default value of `r2_score`. This influences the `score` method of all the multioutput regressors (except for [MultiOutputRegressor](#)).

```
set_params(**params)
```

[\[source\]](#)

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as [Pipeline](#)). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Parameters:

****params : dict**






Estimator parameters.

Returns:

self : estimator instance

or instance.

Examples using `sklearn.linear_model.Ridge`








[Compressive sensing: tomography reconstruction with L1 prior \(Lasso\)](#)

[Prediction Latency](#)

[Comparison of kernel ridge and Gaussian process regression](#)

[Plot Ridge coefficients as a function of the regularization](#)

[Ordinary Least Squares and Ridge Regression Variance](#)



[Plot Ridge coefficients as a function of the L2 regularization](#)

[HuberRegressor vs Ridge on dataset with strong outliers](#)

[Polynomial and Spline interpolation](#)

[Poisson regression and non-normal loss](#)

[Common pitfalls in the interpretation of coefficients of linear models](#)