

# sklearn.linear\_model.Perceptron

```
class sklearn.linear_model.Perceptron(*, penalty=None, alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=1000, tol=0.001, shuffle=True, verbose=0, eta0=1.0, n_jobs=None, random_state=0, early_stopping=False, validation_fraction=0.1, n_iter_no_change=5, class_weight=None, warm_start=False) ¶
```

[\[source\]](#)

Linear perceptron classifier.

Read more in the [User Guide](#).

## Parameters:

**penalty : {'l2','l1','elasticnet'}, default=None**

The penalty (aka regularization term) to be used.

**alpha : float, default=0.0001**

Constant that multiplies the regularization term if regularization is used.

**l1\_ratio : float, default=0.15**

The Elastic Net mixing parameter, with  $0 \leq \text{l1\_ratio} \leq 1$ . `l1_ratio=0` corresponds to L2 penalty, `l1_ratio=1` to L1. Only used if `penalty='elasticnet'`.

*New in version 0.24.*

**fit\_intercept : bool, default=True**

Whether the intercept should be estimated or not. If False, the data is assumed to be already centered.

**max\_iter : int, default=1000**

The maximum number of passes over the training data (aka epochs). It only impacts the behavior in the `fit` method, and not the `partial_fit` method.

*New in version 0.19.*

**tol : float, default=1e-3**

The stopping criterion. If it is not None, the iterations will stop when  $(\text{loss} > \text{previous\_loss} - \text{tol})$ .

*New in version 0.19.*

**shuffle : bool, default=True**

Whether or not the training data should be shuffled after each epoch.

**verbose : int, default=0**

The verbosity level.

**eta0 : float, default=1**

Constant by which the updates are multiplied.

**n\_jobs : int, default=None**

The number of CPUs to use to do the OVA (One Versus All, for multi-class problems) computation. `None` means 1 unless in a `joblib.parallel_backend` context. `-1` means using all processors. See [Glossary](#) for more details.

**random\_state : int, RandomState instance, default=None**

Used to shuffle the training data, when `shuffle` is set to `True`. Pass an int for reproducible output across multiple function calls. See [Glossary](#).

**validation\_fraction : bool, default=False**

Whether to use early stopping to terminate training when validation. score is not improving. If set to True, it will automatically set aside a stratified fraction of training data as validation and terminate training when validation score is not improving by at least tol for n\_iter\_no\_change consecutive epochs.

*New in version 0.20.*

**validation\_fraction : float, default=0.1**

The proportion of training data to set aside as validation set for early stopping. Must be between 0 and 1. Only used if early\_stopping is True.

*New in version 0.20.*

**n\_iter\_no\_change : int, default=5**

Number of iterations with no improvement to wait before early stopping.

*New in version 0.20.*

**class\_weight : dict, {class\_label: weight} or "balanced", default=None**

Preset for the class\_weight fit parameter.

Weights associated with classes. If not given, all classes are supposed to have weight one.

The "balanced" mode uses the values of y to automatically adjust weights inversely proportional to class frequencies in the input data as  $n\_samples / (n\_classes * np.bincount(y))$ .

**warm\_start : bool, default=False**

When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erase the previous solution. See [the Glossary](#).

**Attributes:**

**classes\_ : ndarray of shape (n\_classes,)**

The unique classes labels.

**coef\_ : ndarray of shape (1, n\_features) if n\_classes == 2 else (n\_classes, n\_features)**

Weights assigned to the features.

**intercept\_ : ndarray of shape (1,) if n\_classes == 2 else (n\_classes,)**

Constants in decision function.

**loss\_function\_ : concrete LossFunction**

The function that determines the loss, or difference between the output of the algorithm and the target values.

**n\_features\_in\_ : int**

Number of features seen during [fit](#).

*New in version 0.24.*

**feature\_names\_in\_ : ndarray of shape (n\_features\_in\_,)**

Names of features seen during [fit](#). Defined only when x has feature names that are all strings.

*New in version 1.0.*

**n\_iter\_ : int**

The actual number of iterations to reach the stopping criterion. For multiclass fits, it is the maximum over every binary fit.

**t\_ : int**

Number of weight updates performed during training. Same as  $(n\_iter\_ * n\_samples)$ .

See also:

[sklearn.linear\\_model.SGDClassifier](#)

Linear classifiers (SVM, logistic regression, etc.) with SGD training.

Notes

`Perceptron` is a classification algorithm which shares the same underlying implementation with `SGDClassifier`. In fact, `Perceptron()` is equivalent to `SGDClassifier(loss="perceptron", eta0=1, learning_rate="constant", penalty=None)`.

References

<https://en.wikipedia.org/wiki/Perceptron> and references therein.

Examples

```
>>> from sklearn.datasets import load_digits
>>> from sklearn.linear_model import Perceptron
>>> X, y = load_digits(return_X_y=True)
>>> clf = Perceptron(tol=1e-3, random_state=0)
>>> clf.fit(X, y)
Perceptron()
>>> clf.score(X, y)
0.939...
```

Methods

<a href="#">decision_function(X)</a>	Predict confidence scores for samples.
<a href="#">densify()</a>	Convert coefficient matrix to dense array format.
<a href="#">fit(X, y[, coef_init, intercept_init, ...])</a>	Fit linear model with Stochastic Gradient Descent.
<a href="#">get_params([deep])</a>	Get parameters for this estimator.
<a href="#">partial_fit(X, y[, classes, sample_weight])</a>	Perform one epoch of stochastic gradient descent on given samples.
<a href="#">predict(X)</a>	Predict class labels for samples in X.
<a href="#">score(X, y[, sample_weight])</a>	Return the mean accuracy on the given test data and labels.
<a href="#">set_params(**params)</a>	Set the parameters of this estimator.
<a href="#">sparsify()</a>	Convert coefficient matrix to sparse format.

decision\_function(X)

[source]

Predict confidence scores for samples.

The confidence score for a sample is proportional to the signed distance of that sample to the hyperplane.

Parameters:

**X : {array-like, sparse matrix} of shape (n\_samples, n\_features)**  
The data matrix for which we want to get the confidence scores.

Returns:

**scores : ndarray of shape (n\_samples,) or (n\_samples, n\_classes)**  
Confidence scores per (n\_samples, n\_classes) combination. In the binary case, confidence score for `self.classes_[1]` where >0 means this class would be predicted.

densify()

[source]

Converts the `coef_` member (back) to a `numpy.ndarray`. This is the default format of `coef_` and is required for fitting, so calling this method is only required on models that have previously been sparsified; otherwise, it is a no-op.

**Returns:**

**self**

Fitted estimator.

`fit(X, y, coef_init=None, intercept_init=None, sample_weight=None)`

[\[source\]](#)

Fit linear model with Stochastic Gradient Descent.

**Parameters:**

**X : {array-like, sparse matrix}, shape (n\_samples, n\_features)**

Training data.

**y : ndarray of shape (n\_samples,)**

Target values.

**coef\_init : ndarray of shape (n\_classes, n\_features), default=None**

The initial coefficients to warm-start the optimization.

**intercept\_init : ndarray of shape (n\_classes,), default=None**

The initial intercept to warm-start the optimization.

**sample\_weight : array-like, shape (n\_samples,), default=None**

Weights applied to individual samples. If not provided, uniform weights are assumed. These weights will be multiplied with `class_weight` (passed through the constructor) if `class_weight` is specified.

**Returns:**

**self : object**

Returns an instance of self.

`get_params(deep=True)`

[\[source\]](#)

Get parameters for this estimator.

**Parameters:**

**deep : bool, default=True**

If True, will return the parameters for this estimator and contained subobjects that are estimators.

**Returns:**

**params : dict**

Parameter names mapped to their values.

`partial_fit(X, y, classes=None, sample_weight=None)`

[\[source\]](#)

Perform one epoch of stochastic gradient descent on given samples.

Internally, this method uses `max_iter = 1`. Therefore, it is not guaranteed that a minimum of the cost function is reached after calling it once. Matters such as objective convergence, early stopping, and learning rate adjustments should be handled by the user.

**Parameters:**

***X : {array-like, sparse matrix}, shape (n\_samples, n\_features)***

Subset of the training data.

***y : ndarray of shape (n\_samples,)***

Subset of the target values.

***classes : ndarray of shape (n\_classes,), default=None***

Classes across all calls to `partial_fit`. Can be obtained by via `np.unique(y_all)`, where `y_all` is the target vector of the entire dataset. This argument is required for the first call to `partial_fit` and can be omitted in the subsequent calls. Note that `y` doesn't need to contain all labels in `classes`.

***sample\_weight : array-like, shape (n\_samples,), default=None***

Weights applied to individual samples. If not provided, uniform weights are assumed.

**Returns:**

***self : object***

Returns an instance of `self`.

`predict(X)`

[\[source\]](#)

Predict class labels for samples in `X`.

**Parameters:**

***X : {array-like, sparse matrix} of shape (n\_samples, n\_features)***

The data matrix for which we want to get the predictions.

**Returns:**

***y\_pred : ndarray of shape (n\_samples,)***

Vector containing the class labels for each sample.

`score(X, y, sample_weight=None)`

[\[source\]](#)

Return the mean accuracy on the given test data and labels.

In multi-label classification, this is the subset accuracy which is a harsh metric since you require for each sample that each label set be correctly predicted.

**Parameters:**

***X : array-like of shape (n\_samples, n\_features)***

Test samples.

***y : array-like of shape (n\_samples,) or (n\_samples, n\_outputs)***

True labels for `x`.

***sample\_weight : array-like of shape (n\_samples,), default=None***

Sample weights.

**Returns:**

***score : float***

Mean accuracy of `self.predict(X)` wrt. `y`.

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as [Pipeline](#)). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Parameters:

**\*\*params : dict**  
Estimator parameters.

Returns:

**self : estimator instance**  
Estimator instance.

Convert coefficient matrix to sparse format.

Converts the `coef_` member to a `scipy.sparse` matrix, which for L1-regularized models can be much more memory- and storage-efficient than the usual `numpy.ndarray` representation.

The `intercept_` member is not converted.

Returns:

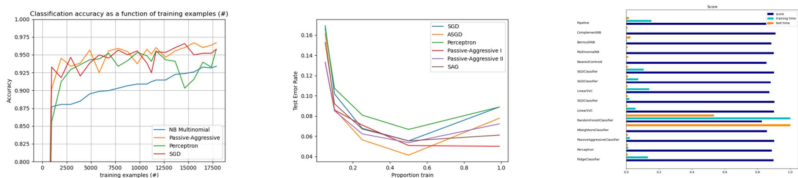
**self**  
Fitted estimator.

Notes

For non-sparse models, i.e. when there are not many zeros in `coef_`, this may actually *increase* memory usage, so use this method with care. A rule of thumb is that the number of zero elements, which can be computed with `(coef_ == 0).sum()`, must be more than 50% for this to provide significant benefits.

After calling this method, further fitting with the `partial_fit` method (if any) will not work until you call `densify`.

Examples using `sklearn.linear_model.Perceptron`



[Out-of-core classification of text documents](#)

[Comparing various online solvers](#)

[Classification of text documents using sparse features](#)