

sklearn.linear_model.RidgeCV

```
class sklearn.linear_model.RidgeCV(alphas=(0.1, 1.0, 10.0), *, fit_intercept=True, normalize='deprecated', scoring=None, cv=None, gcv_mode=None, store_cv_values=False, alpha_per_target=False)
```

[\[source\]](#)

Ridge regression with built-in cross-validation.

See glossary entry for [cross-validation estimator](#).

By default, it performs efficient Leave-One-Out Cross-Validation.

Read more in the [User Guide](#).

Parameters:

alphas : *ndarray of shape (n_alphas,)*, **default=(0.1, 1.0, 10.0)**

Array of alpha values to try. Regularization strength; must be a positive float. Regularization improves the conditioning of the problem and reduces the variance of the estimates. Larger values specify stronger regularization. Alpha corresponds to $1 / (2C)$ in other linear models such as [LogisticRegression](#) or [LinearSVC](#). If using Leave-One-Out cross-validation, alphas must be positive.

fit_intercept : *bool*, **default=True**

Whether to calculate the intercept for this model. If set to false, no intercept will be used in calculations (i.e. data is expected to be centered).

normalize : *bool*, **default=False**

This parameter is ignored when `fit_intercept` is set to False. If True, the regressors X will be normalized before regression by subtracting the mean and dividing by the l2-norm. If you wish to standardize, please use [StandardScaler](#) before calling `fit` on an estimator with `normalize=False`.

Deprecated since version 1.0: `normalize` was deprecated in version 1.0 and will be removed in 1.2.

scoring : *str, callable*, **default=None**

A string (see model evaluation documentation) or a scorer callable object / function with signature `scorer(estimator, X, y)`. If None, the negative mean squared error if cv is 'auto' or None (i.e. when using leave-one-out cross-validation), and r2 score otherwise.

cv : *int, cross-validation generator or an iterable*, **default=None**

Determines the cross-validation splitting strategy. Possible inputs for cv are:

- None, to use the efficient Leave-One-Out cross-validation
- integer, to specify the number of folds.
- [CV splitter](#),
- An iterable yielding (train, test) splits as arrays of indices.

For integer/None inputs, if `y` is binary or multiclass, [StratifiedKFold](#) is used, else, [KFold](#) is used.

Refer [User Guide](#) for the various cross-validation strategies that can be used here.

gcv_mode : *{'auto', 'svd', 'eigen'}*, **default='auto'**

Flag indicating which strategy to use when performing Leave-One-Out Cross-Validation. Options are:

```
'auto' : use 'svd' if n_samples > n_features, otherwise use 'eigen'
'svd'  : force use of singular value decomposition of X when X is
        dense, eigenvalue decomposition of X^T.X when X is sparse.
'eigen': force computation via eigendecomposition of X.X^T
```

The 'auto' mode is the default and is intended to pick the cheaper option of the two depending on the shape of the training data.

store_cv_values : bool, default=False

Flag indicating if the cross-validation values corresponding to each alpha should be stored in the `cv_values_` attribute (see below). This flag is only compatible with `cv=None` (i.e. using Leave-One-Out Cross-Validation).

alpha_per_target : bool, default=False

Flag indicating whether to optimize the alpha value (picked from the `alphas` parameter list) for each target separately (for multi-output settings: multiple prediction targets). When set to `True`, after fitting, the `alpha_` attribute will contain a value for each target. When set to `False`, a single alpha is used for all targets.

New in version 0.24.

Attributes:

cv_values_ : ndarray of shape (n_samples, n_alphas) or shape (n_samples, n_targets, n_alphas), optional

Cross-validation values for each alpha (only available if `store_cv_values=True` and `cv=None`). After `fit()` has been called, this attribute will contain the mean squared errors if `scoring is None` otherwise it will contain standardized per point prediction values.

coef_ : ndarray of shape (n_features) or (n_targets, n_features)

Weight vector(s).

intercept_ : float or ndarray of shape (n_targets,)

Independent term in decision function. Set to 0.0 if `fit_intercept = False`.

alpha_ : float or ndarray of shape (n_targets,)

Estimated regularization parameter, or, if `alpha_per_target=True`, the estimated regularization parameter for each target.

best_score_ : float or ndarray of shape (n_targets,)

Score of base estimator with best alpha, or, if `alpha_per_target=True`, a score for each target.

New in version 0.23.

n_features_in_ : int

Number of features seen during [fit](#).

New in version 0.24.

feature_names_in_ : ndarray of shape (n_features_in_,)

Names of features seen during [fit](#). Defined only when `x` has feature names that are all strings.

New in version 1.0.

See also:

[Ridge](#)

Ridge regression.

[RidgeClassifier](#)

Classifier based on ridge regression on `{-1, 1}` labels.

[RidgeClassifierCV](#)

Ridge classifier with built-in cross validation.

Examples

```
>>> from sklearn.datasets import load_diabetes
>>> from sklearn.linear_model import RidgeCV
>>> X, y = load_diabetes(return_X_y=True)
>>> clf = RidgeCV(alphas=[1e-3, 1e-2, 1e-1, 1]).fit(X, y)
>>> clf.score(X, y)
0.5166...
```

Methods

| | |
|---|--|
| fit (X, y[, sample_weight]) | Fit Ridge regression model with cv. |
| get_params ([deep]) | Get parameters for this estimator. |
| predict (X) | Predict using the linear model. |
| score (X, y[, sample_weight]) | Return the coefficient of determination of the prediction. |
| set_params (**params) | Set the parameters of this estimator. |

fit(X, y, sample_weight=None)

[source]

Fit Ridge regression model with cv.

Parameters:

- X : ndarray of shape (n_samples, n_features)**
Training data. If using GCV, will be cast to float64 if necessary.
- y : ndarray of shape (n_samples,) or (n_samples, n_targets)**
Target values. Will be cast to X's dtype if necessary.
- sample_weight : float or ndarray of shape (n_samples,), default=None**
Individual weights for each sample. If given a float, every sample will have the same weight.

Returns:

- self : object**
Fitted estimator.

Notes

When sample_weight is provided, the selected hyperparameter may depend on whether we use leave-one-out cross-validation (cv=None or cv='auto') or another form of cross-validation, because only leave-one-out cross-validation takes the sample weights into account when computing the validation score.

get_params(deep=True)

[source]

Get parameters for this estimator.

Parameters:

- deep : bool, default=True**
If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns:

- params : dict**
Parameter names mapped to their values.

Predict using the linear model.

Parameters:

X : array-like or sparse matrix, shape (n_samples, n_features)

Samples.

Returns:

C : array, shape (n_samples,)

Returns predicted values.

score(X, y, sample_weight=None)

[source]

Return the coefficient of determination of the prediction.

The coefficient of determination R^2 is defined as $(1 - \frac{u}{v})$, where u is the residual sum of squares $((y_{\text{true}} - y_{\text{pred}})** 2).sum()$ and v is the total sum of squares $((y_{\text{true}} - y_{\text{true}.mean()}) ** 2).sum()$. The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of y , disregarding the input features, would get a R^2 score of 0.0.

Parameters:

X : array-like of shape (n_samples, n_features)

Test samples. For some estimators this may be a precomputed kernel matrix or a list of generic objects instead with shape $(n_{\text{samples}}, n_{\text{samples_fitted}})$, where $n_{\text{samples_fitted}}$ is the number of samples used in the fitting for the estimator.

y : array-like of shape (n_samples,) or (n_samples, n_outputs)

True values for X .

sample_weight : array-like of shape (n_samples,), default=None

Sample weights.

Returns:

score : float

R^2 of `self.predict(X)` wrt. y .

Notes

The R^2 score used when calling `score` on a regressor uses `multioutput='uniform_average'` from version 0.23 to keep consistent with default value of `r2_score`. This influences the `score` method of all the multioutput regressors (except for `MultiOutputRegressor`).

set_params(**params)

[source]

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as `Pipeline`). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

Parameters:

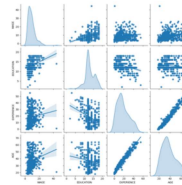
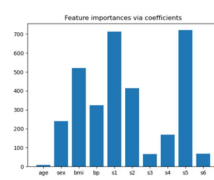
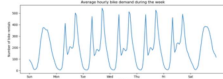
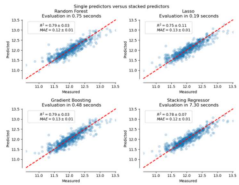
****params : dict**

Estimator parameters.

Returns:

self : estimator instance

Examples using `sklearn.linear_model.RidgeCV`



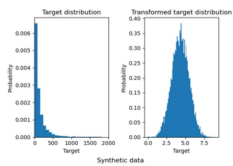
[Combine predictors using stacking](#)

[Time-related feature engineering](#)

[Model-based and sequential feature selection](#)

[Common pitfalls in the interpretation of coefficients of linear models](#)

[Face completion with a multi-output estimators](#)



[Effect of transforming the targets in regression model](#)