

Computer Networks

Data Link Layer

Paolo Costa
costa@cs.vu.nl
<http://www.cs.vu.nl/~costa>
Vrije Universiteit Amsterdam

(Version May 9, 2008)

Paolo Costa

03 - Data Link Layer

1 / 55

Data Link Layer

- **Physical:** Describes the transmission of raw bits in terms of mechanical and electrical issues.
- **Data Link:** Describes how a shared communication channel can be accessed, and [how a data frame can be reliably transmitted](#).
- **Network:** Describes how routing is to be done. Mostly needed in subnets.
- **Transport:** The hardest one: generally offers connection-oriented as well as connectionless services, and varying degrees of reliability. This layer provides the actual network interface to [applications](#).
- **Application:** Contains the stuff that users see: e-mail, remote logins, the Web's exchange protocol, etc.

Note


We'll just concentrate on transmission issues. Channel access is discussed in Chapter 4

Paolo Costa

03 - Data Link Layer

Introduction 2 / 55

Design Issues



- 1 Provide well-defined interface to network layer
- 2 Handle transmission errors
- 3 Regulate flow of data: get sender and receiver in the same pace

Paolo Costa

03 - Data Link Layer

Design Issues 3 / 55

Basic Services

- Network layer passes a number of bits (**frame**) to the data link layer.
- Data link layer is responsible for transmitting the frame to the destination machine.
- Receiving layer passes received frame to its network layer.

Basic services commonly provided:

- unacknowledged connectionless service (LANs)
- acknowledged connectionless service (Wireless systems).
- acknowledged connection-oriented service (WANs).

Question

- Why are we so concerned about error-free frame transmissions? Can't the higher layers take care of that?
It's not mandatory but it may improve efficiency (fine-grained recovery: frames vs. messages)

Paolo Costa

03 - Data Link Layer

Design Issues 4 / 55


Notes

Notes

Notes

Notes

Transmission



- The actual transmissions follow the path of (b) but it is easier to think in terms of two data link process communicating using a data link protocol (a)


Paolo Costa

03 - Data Link Layer

Design Issues 5 / 55

Notes

Routing



- The network layer consists of routing
 - they are connected through point-to-point links
- The router would really its packet to be sent correctly, guaranteed, and in the order it was issued.
- It is up to the data link to make unreliable connections look perfect, or at least, fairly good

Paolo Costa

03 - Data Link Layer


Design Issues 6 / 55

Notes

Frames

Counting

- The physical layer doesn't do much: it just pumps bits from one end to the other.
- But things may go wrong
 - ⇒ the data link layer needs a means to do retransmissions.
- The unit of retransmission is a **frame** (which is just a fixed number of bits).
- Problem:** How can we break up a bit stream into frames?
 - naive solution: **counting**



(a) Without errors (b) With errors

Paolo Costa

03 - Data Link Layer

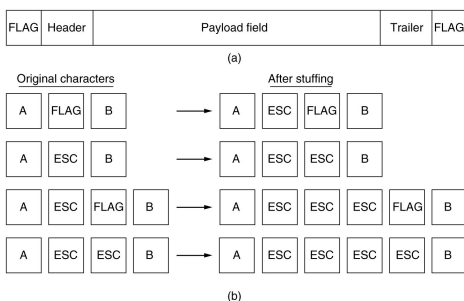
Design Issues 7 / 55

Notes

Frames

Byte Stuffing

- Byte stuffing:** Mark the beginning and end of a *byte* frame with two special **flag bytes** – a special bit sequence (e.g. 01111110). If such bytes appear in the original frame, escape them:



Paolo Costa

03 - Data Link Layer

Design Issues 8 / 55

Notes

Frames

Bit Stuffing

- Byte stuffing is closely tied to the use of 8-bit character
- Solution**
 - ⇒ **Bit stuffing**: Escape the flag byte (e.g., 01111110) through an additional bit
 - whenever the sender's data link encounters five consecutive 1s in the data, it automatically stuffs a 0 bit

(a) 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0

(b) 0 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 1 0 0 1 0

↑ ↑ ↑
Stuffed bits

(c) 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0

(a) frame to send (b) frame transmitted over the wire (c) reconstructed frame

Notes

Error Correction and Detection

- Problem**: Suppose something went wrong during frame transmission. How do we actually **notice** that something's wrong, and can it be corrected **by the receiver**?
- Definition**: The Hamming distance between two frames \vec{a} and \vec{b} is the number of bits at the same position that differ.

Example

10001001 and 10110001 are at Hamming distance **3**

1	0	0	0	1	0	0	1
1	0	1	1	0	0	0	1
							⊕ ← XOR
0	0	1	1	1	0	0	0

- To **detect** d errors, you need a distance $d + 1$ code
 - d -single bit errors cannot change a valid codeword in another
- To **correct** d errors, you need a distance $2d + 1$ code
 - this way, even with $2d$ changes, the original codeword is still closer than any other codeword

Notes

Error Detection

Parity

- Add a bit to a bit string such that the total number of 1-bits is even (or odd)
 - e.g., (**even parity**) 1011010 ⇒ 1011010**0**
 - (**odd parity**) 1011010 ⇒ 1011010**1**
- The distance between two (**legal**) frames is at least $d = 2$.
 - any single-bit error produces a code word with the wrong parity
- We can **detect** a **single** error (i.e., $k = 1$ and $d = 2$)
 - remember that to **detect** d errors, you need a distance $d + 1$ code

Notes

Error Correction

Example

- Consider a code with only four valid codewords:
 - 0000000000
 - 0000011111
 - 1111100000
 - 1111111111
- This code has distance **5**
 - The **XOR** between any pair of the above codeword gives at least five 1 bits
 - ⇒ it means it can correct double errors
 - remember that to **correct** d errors, you need a distance $2d + 1$ code
- E.g., if the codeword 0000000111 is received, the receiver knows that the original must have been 0000001111
- If, however, a **triple** error changes 0000000000 into 0000000111, the error will not be corrected properly

Notes

Error Correction

Lower Bound

- **Problem:** we want to design a code with m message bits and r check bits to correct all single errors
- Each of the 2^m legal messages has n illegal codewords at distance 1
 - these are obtained by systematically change each of the n bits in the n -bit codeword
- Each of the 2^m legal messages require $n + 1$ bit patterns dedicated to it
- The total number of bit pattern is 2^n
 $\Rightarrow (n + 1)2^m \leq 2^n \quad n = \frac{m}{1-r} \quad (m + r + 1) \leq 2^r$
- Given m , it is possible to have a lower bound to the number r of check bits needed to correct **single** errors

Paolo Costa

03 - Data Link Layer

Error Control 13 / 55

Error Correction: Hamming

Definition

- The theoretical lower limit can be achieved using a method due to **Hamming** (1950)
 - Enumerate all bits in a codeword starting with bit 1 at the left
 - The bits at position 2^k ($k \geq 0$) are check bits, the rest (3,5,6,7,...) are filled with the m data bits
 - Every check bit is used as a parity bit for those positions to which it contributes
- | | b1 | b2 | b3 | b4 | b5 | b6 | b7 | b8 | b9 | b10 | b11 |
|---|----|----|----|----|----|----|----|----|----|-----|-----|
| 1 | X | | X | | X | | X | | X | | X |
| 2 | | X | X | | | X | X | | | X | X |
| 4 | | | | X | X | X | X | | | | |
| 8 | | | | | | | | X | X | X | X |
- Rewrite every data bit as sums of power of 2
 - e.g., $7 = 1 + 2 + 4$ and $11 = 1 + 2 + 8$
 - A bit is checked by just those bits occurring in its expansion
 - e.g., bit 11 is checked by bit 1, 2, and 8

Paolo Costa

03 - Data Link Layer

Error Control 14 / 55

Error Correction: Hamming

Coding

	b1	b2	b3	b4	b5	b6	b7	b8	b9	b10	b11
1	X		X		X		X		X		X
2		X	X			X	X			X	X
4				X	X	X	X				
8								X	X	X	X
	1	0	1	1	1	0	0	1	0	0	1

- Check bits are used to even out the number of 1 bits they contribute to
 - e.g., assume the message is **1100001**
 - Check bit at position 2, is used to even out the bits for positions 2, 3, 6, 7, 10, and 11
 $\Rightarrow b2 = 0$ as the number of 1s is even (b3 and b11)
 - Other check bits are computed in the same way
 - The final codeword is **1011100101**

Paolo Costa

03 - Data Link Layer

Error Control 15 / 55

Error Correction: Hamming

Decoding

- If a check bit at position p is wrong upon receipt, the receiver increments a counter v with p
- The value of v will, in the end, give the position of the wrong bit, which should then be swapped. E.g.:

	b1	b2	b3	b4	b5	b6	b7	b8	b9	b10	b11
1	X		X		X		X		X		X
2		X	X			X	X			X	X
4				X	X	X	X				
8								X	X	X	X
S:	1	0	1	1	1	0	0	1	0	0	1
R:	1	0	1	1	1	0	0	1	1	0	1
C:	0	0	1	1	1	0	0	0	1	0	1
F:	1	0	1	1	1	0	0	1	0	0	1

S: string sent

R: string received

C: string corrected on the check bits: #1 and #8 corrected \Rightarrow bit #9 is wrong ($v = 1 + 8$)

F: final result after correction

Paolo Costa

03 - Data Link Layer

Error Control 16 / 55

Notes

Notes

Notes

Notes

Error Correction: Hamming

Burst Errors

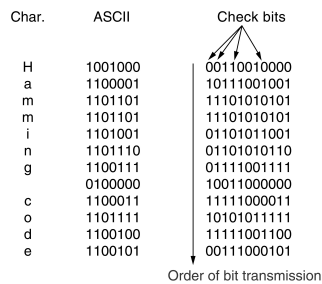
- Hamming codes can only detect **single** errors
- However, there is a trick to correct burst errors

- A sequence of k codewords are arranged as a matrix, one codeword per **row**

- Data are transmitted one **column** at a time

- If a burst error of length k occurs, at most **1** bit per codeword is affected

- Hamming code can correct one error per codeword
- ⇒ the entire block can be restored



Paolo Costa

03 - Data Link Layer

Error Control

17 / 55

Error Detection

CRC

- Error correcting codes are simply too expensive
 - e.g., for 10,000 bit blocks, 10 check bits are needed
 - ⇒ only use error detection combined with retransmissions
- Cyclic Redundancy Check (CRC)**, a.k.a. **polynomial** code, is based upon treating bit strings as polynomials with coefficient of **0** and **1** only
 - a k -bit frame is regarded as the coefficient list for a polynomial with k terms, from x^{k-1} to x^0
 - e.g., a 6-bit block

$$M(x) = m_{k-1}x^{k-1} + \dots + m_1x^1 + m_0x^0$$

$$\begin{aligned} 110001 &\Rightarrow 1\dot{x}^5 + 1\dot{x}^4 + 0\dot{x}^3 + 0\dot{x}^2 + 0\dot{x}^1 + 1\dot{x}^0 \\ &= x^5 + x^4 + 1 \end{aligned}$$

- Polynomial arithmetic is done in modulo 2, according to the rule of algebraic field theory

Paolo Costa

03 - Data Link Layer

Error Control

18 / 55

Polynomial Arithmetic

Addition and Subtraction

- Since the coefficients are constrained to a single bit, any math operation on CRC polynomials must map the coefficients of the result to either **0** or **1**
- Addition**

$$(x^3 + x) + (x + 1) = x^3 + 2x + 1 \equiv x^3 + 1$$

- note that $2x$ becomes zero because addition of coefficient is done modulo 2

$$2x = x + x = (1 + 1)x = 0x = 0$$

- analogous to **exclusive OR (XOR)**

$$1010 \oplus 0011 \rightarrow 1001$$

- Subtraction**
 - works like addition (only absolute values are used)

$$(x^4 + x^2 + 1) - (x + 1) = x^4 + x^2 - x \equiv x^4 + x^2 + x$$

Paolo Costa

03 - Data Link Layer

Error Control

19 / 55

Polynomial Arithmetic

Multiplication and Division

- Multiplication**
 - It follows the rule of traditional polynomials but, again, addition of coefficient is done module 2

$$(x^2 + x)(x + 1) = x^3 + x^2 + x^2 + 1 = x^3 + 2x^2 + 1 \equiv x^3 + 1$$

- Division**
 - we can also divide polynomials mod 2 and find the quotient and remainder

$$\frac{x^3 + x^2 + x}{x + 1} = (x^2 + 1) - \frac{1}{x - 1}$$

- or in other words:

$$(x^3 + x^2 + x) = (x^2 + 1)(x + 1) - 1$$

Paolo Costa

03 - Data Link Layer

Error Control

20 / 55

Notes

Notes

Notes

Notes

CRC

Algorithm

- Let's go back our problem
- Sender and receiver agree upon a **generator polynomial** $G(x)$
 - both the high- and low-order bits of the generator must be 1
 - the frame $M(x)$ must be longer than the generator ($m > r$)
- Idea**: append a checksum to the end of the frame in such a way that the polynomial obtained is divisible (i.e., no remainder) by $G(x)$
 - when the receivers gets the checksummed frame, it tries dividing it by $G(x)$
 - if there is a remainder, there has been a transmission error
- Algorithm**
 - append r zero bits to the low order of the frame
 - now it contains $m + r$ bits and corresponds to $x^r M(x)$
 - divide $x^r M(x)$ by $G(x)$
 - subtract (mod 2) the remainder from $x^r M(x)$. The result is the checksummed frame $T(x)$ to be transmitted

Paolo Costa

03 - Data Link Layer

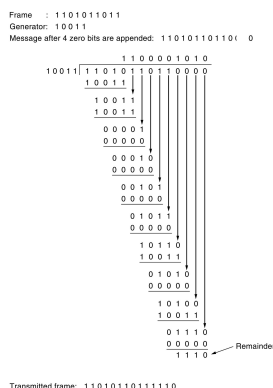
Error Control 21 / 55

Notes

CRC

Division Example

- Check the input bit above the leftmost divisor
 - If it is 0, do nothing
 - If it is 1, the divisor is **XORed** into (i.e., subtracted from) the input
- Move the divisor to the right by one bit
- The process is repeated until the divisor reaches the right-hand end of the input row
- The left bits are the remainder of the division to be appended to the frame



Paolo Costa

03 - Data Link Layer

Error Control 22 / 55

Notes

CRC

Decoding

- If a transmission error occurs, instead of the original polynomial $T(x)$, $T(x) + E(x)$ arrives
 - each 1 bit in $E(x)$ corresponds to a bit that has been inverted
 - if there are k 1 bits in $E(x)$, k single bit errors have occurred
 - $E(x)$ is characterized by an initial 1, a mixture of 0s and 1s, and a final 1
- The receiver divided the received frame by $G(x)$, i.e., $[T(x) + E(x)]/G(x)$ and computes the remainder
 - in case of error, the remainder will be $E(x)/G(x)$
- If there is a **single-bit** error, $E(x) = x^i$ (being i the position of the wrong bit)
 - if $G(x)$ contains two or more terms, it will **never** divide $E(x) \Rightarrow$ all single-bit errors will be detected
 - the simplest error-detection system, the parity bit, is in fact a trivial CRC: it uses the two-bit-long divisor 11

Paolo Costa

03 - Data Link Layer

Error Control 23 / 55

Notes

CRC

Burst Error

- By carefully choosing $G(x)$, **burst** errors can be correctly detected
 - a burst error is a contiguous sequence of bits, containing wrong bits
- It can be shown, that a polynomial code with r check bits, will detect all bursts errors of length $\leq r$
- If the burst length is $r + 1$, the probability of accepting an incorrect frame is $\frac{1}{2^{r-1}}$
- It can also be shown that for bursts longer than $r + 1$ bit, such probability becomes $\frac{1}{2^r}$, assuming that all bit patterns are equally likely (not always true)
- Certain polynomial have become standards
 - e.g., the one used in IEEE 802 (Wi-Fi) is:
$$x^{32} + x^{26} + x^{23} + x^{22} + x^{16} + x^{12} + x^{11} + x^{10} + x^8 + x^7 + x^5 + x^4 + x^2 + x^1 + 1$$
 - it detects all bursts of length 32 or less and all bursts affecting an **odd** number of bits

Paolo Costa

03 - Data Link Layer

Error Control 24 / 55

Notes

CRC

Applications

- CRC is not used by the data link layer but also by applications to check the correctness of a given file
- e.g., WinZip



Data Link Layer Protocols

- Concentrated on design aspects and error control
 - frame size (stuffing)
 - error correction (Hamming Code)
 - error detection (CRC)
- Now: basic protocols and real-world examples

Some basic assumptions:

- We have a machine *A* that wants to send data to machine *B*
- There is always enough data for *A* to send
- There is a well-defined interface to the network layer, and to the physical layer
 - void from_network_layer(packet *)
 - void to_network_layer(packet *)
 - void from_physical_layer(packet *)
 - void to_physical_layer(packet *)
- The receiver generally waits for an event to happen by calling
 - void wait_for_event(event_type *event)

Unrestricted Simplex Protocol

```
1 typedef enum {false, true} boolean;
2 typedef unsigned int seq_nr;
3 typedef struct {unsigned char data[MAX_PKT];} packet;
4 typedef enum {data, ack, nak} frame_kind;
5
6 typedef struct {
7     frame_kind kind; /* what kind of a frame is it? */
8     seq_nr seq; /* sequence number */
9     seq_nr ack; /* acknowledgment number */
10    packet info; /* the network layer packet */
11 } frame;
12
13 typedef enum {frame_arrival} event_type;
14
15 void sender1(void) {
16     frame s; packet buffer;
17     while (true) {
18         from_network_layer(&buffer);
19         s.info = buffer;
20         to_physical_layer(&s);
21     }
22 }
23
24 void receiver1(void) {
25     frame r; event_type event;
26     while (true) {
27         wait_for_event(&event);
28         from_physical_layer(&r);
29         to_network_layer(&r.info);
30     }
31 }
```

Unrestricted Simplex Protocol

Issues

Question

- What are some of the underlying assumptions here? How does the flow control manifest itself?

Underlying assumptions are error-free transmission, and sender-receiver are in pace.

- **Issue:** Fast senders would make receivers collapse

Question

- Would a simple delay on sender's side fix the problem ?

Not really! It's hard to predict when receiver's congestion may occur.

- Assuming always the worst-case behavior would be highly inefficient

Notes

Notes

Notes

Notes

Simplex Stop-and-Wait

```
1 typedef enum (frame_arrival) event_type;
2 #include "protocol.h"
3
4 void sender2(void) {
5     frame s; packet buffer;
6     event_type event;
7     while (true) {
8         from_network_layer(&buffer);
9         s.info = buffer;
10        to_physical_layer(&s);
11        wait_for_event(&event);
12    }
13 }
14
15 void receiver2(void) {
16     frame r, s; event_type event;
17     while (true) {
18         wait_for_event(&event);
19         from_physical_layer(&r);
20         to_network_layer(&r.info);
21         to_physical_layer(&s);
22     }
23 }
```

- Here the receiver provide the feedback to the sender
 - it sends a little dummy packet (**acknowledgment**) back to the sender
- After having sent a frame, the sender is required to wait for the ack

Simplex Protocol for Noisy Channel

Question

- What are the assumptions in this case? **We are assuming only error-free transmission.**

- Ok, so assume that damaged frames can be detected

Question

- Idea: send the ack only upon correct receipt! Works ? **No, because also the ack can be lost.**

Question

- Second idea: let the sender use a timer by which it simply retransmits unacknowledged frames after some time! Works ? **No again! Because the data link layer cannot detect duplicate transmissions.**
- Let's use sequence numbers!
 - **Problem:** sequence number cannot go on for ever and they wast space
 - **Solution:** no worries, we need just two (0 & 1)

Simplex Protocol for Noisy Channel

Sender

- **next_frame_to_send** indicates which is the last message sent
- A timer is started after sending a message
- Only if the correct ack is received, the next message is sent

```
1 #define MAX_SEQ 1
2 typedef enum (frame_arrival, cksum_err, timeout) event_type;
3 void sender3(void) {
4     int next_frame_to_send;
5     frame s;
6     packet buffer;
7     event_type event;
8
9     next_frame_to_send = 0;
10    from_network_layer(&buffer);
11    while (true) {
12        s.info = buffer;
13        s.seq = next_frame_to_send;
14        to_physical_layer(&s);
15        start_timer(s.seq);
16        wait_for_event(&event);
17
18        if (event == frame_arrival) {
19            from_physical_layer(&s);
20            if (s.ack == next_frame_to_send) {
21                stop_timer(s.ack);
22                from_network_layer(&buffer);
23                inc(next_frame_to_send);
24            } } }
```

Simplex Protocol for Noisy Channel

Receiver

- **frame_expected** enables filtering out duplicate
- Its value is increased each time a frame is delivered to the network layer
 - ⇒ only **two** values (0 and 1) are needed

```
1 void receiver3(void) {
2     seq_nr frame_expected; frame r, s; event_type event;
3
4     frame_expected = 0;
5     while (true) {
6         wait_for_event(&event);
7         if (event == frame_arrival) {
8             from_physical_layer(&r);
9             if (r.seq == frame_expected) {
10                to_network_layer(&r.info);
11                inc(frame_expected);
12            }
13            s.ack = 1 - frame_expected;
14            to_physical_layer(&s);
15        }
16    }
17 }
```

Notes

Notes

Notes

Notes

From Simplex to Duplex

- **Problem:** We want to allow symmetric frame transmission between two communicating parties, rather than transmission in one direction.
 - don't waste channels, so use the same channel for **duplex** communication.
- **Solution:** Just transmit frames, but distinguish between data, acknowledgments (acks), and possibly negative acks (nacks) in the frame's type field.
- **Idea:** If the other party is going to send data as well, it might as well send acknowledgments along with its data frames
⇒ piggybacking.

Question

- What's good and bad about piggybacking?
Good: save bandwidth.
Bad: poor performance with irregular transmission rate.

Paolo Costa

03 - Data Link Layer

Basic Protocols 33 / 55

Sliding Windows

- **Principle:** Rather than just sending a single frame at a time, permit the sender to transmit a set of frames, called the **sending window**.
 - a frame is removed from the sending window iff it has been acknowledged.
 - the receiver has a **receiving window** containing frames it is permitted to receive.
- A damaged frame is kept in the receiving window until a correct version is received. Also, frame **N** is kept in the window until frame **N-1** has been received.
 - if both window sizes equal one, we are dealing with the stop-and-wait protocols.

Question


- Is there a relationship between **window size**, **transmission rate**, and **propagation delay**? **If the propagation delay is high and the window size is small, transmission rate decreases dramatically because the sender while waiting for acks cannot send new messages**

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols 34 / 55

Window Size & Sequence Number



- (a) Initially
(b) After sending frame #1.
(c) After receiving frame #1.
(d) After receiving ack for frame #1

Question

- How would you interpret the shaded areas?
Shaded area is to-be-acknowledged (at sender) and to-be-received next (at receiver).

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols 35 / 55

1-Bit Sliding Window

```
1 void protocol4(void) {
2     seq_nr next_frame_to_send, frame_expected;
3     frame r, s; packet buffer;
4     event_type event;
5
6     next_frame_to_send = 0; frame_expected = 0;
7     from_network_layer(s,buffer);
8     s.info = buffer;
9     s.seq = next_frame_to_send;
10    s.ack = 1 - frame_expected;
11    to_physical_layer(s); start_timer(s.seq);
12
13    while (true) {
14        wait_for_event(&event);
15        if (event == frame_arrival) {
16            from_physical_layer(&r);
17            if (r.seq == frame_expected) {
18                to_network_layer(r,info);
19                inc(frame_expected);
20            }
21            if (r.ack == next_frame_to_send) {
22                from_network_layer(s,buffer);
23                inc(next_frame_to_send);
24            }
25        }
26        s.info = buffer;
27        s.seq = next_frame_to_send;
28        s.ack = 1 - frame_expected;
29        to_physical_layer(s); start_timer(s.seq);
30    }
31 }
```

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols 36 / 55

- **next_frame_to_send** tells which frame the sender is trying to send
- **frame_expected** tells which frames the receiver is expecting
- If the **received frame** is the one expected, the frame is passed to the network layer
- If the **received ack** is the one expected, the next frame is sent

Notes

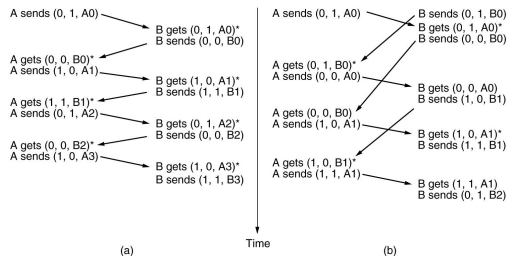
Notes

Notes

Notes

1-Bit Sliding Window

Simultaneous Transmissions



- All things go well, but behavior is a bit strange when *A* and *B* transmit simultaneously.
- We are transmitting more than once, just because the two senders are more or less out of sync.

Notes

Error Control



Problem: What should the receiver do if a frame is damaged ?

- **Go back n**
Simply request retransmission of all frames starting from frame #N. If any other frames had been received in the meantime (and stored in the receiver's window), they'll just be ignored
- **Selective Repeat**
Request just retransmission of the damaged frame, and wait until it comes in before delivering any frames after that

Notes

Error Control

Example



(a) Go back n (b) Selective Repeat.

Notes

Error Control

Pros & Cons

- **Go-back-N** is really simple: the sender keeps a frame in its window until it is acknowledged.
 - if the window is full, the network layer is not allowed to submit new packets.
 - the receiver hardly needs to keep an account on what happens: if a frame is damaged, its successors in the receive window are ignored.
 - it is equivalent to a **receive window of size 1**
- **Selective repeat** seems to do better because frames aren't discarded, but the administration is much harder
 - it also requires large amounts of memory if the window is large
- Trade-off between **bandwidth** and data link layer **buffer space**

Notes

Go-back-n (1/2)

```
1  #define MAX_SEQ 7 /* should be 2^n - 1 */
2  typedef enum {frame_arrival, cksum_err,
3  timeout, network_layer_ready} event_type;
4  #include "protocol.h"
5
6  static boolean between(seq_nr a, seq_nr b, seq_nr c) {
7      /* Return TRUE iff a <= b < c (cyclic) */
8      ...
9  }
10
11 static void send_data(
12     seq_nr frame_nr, seq_nr frame_expected, packet buffer[]) {
13     frame s;
14     s.info = buffer[frame_nr]; s.seq = frame_nr;
15     s.ack = (frame_expected + MAX_SEQ) % (MAX_SEQ + 1);
16     to_physical_layer(s); start_timer(frame_nr);
17 }
18
19 void protocol5(void) {
20     seq_nr next_frame_to_send, ack_expected, frame_expected;
21     frame r;
22     packet buffer[MAX_SEQ + 1];
23     seq_nr nbuffered, i;
24     event_type event;
25
26     enable_network_layer();
27     ack_expected = 0; next_frame_to_send = 0; frame_expected = 0;
28     nbuffered = 0;
```

- `enable_network_layer` prevents network layer to input more frames when there are `MAX_SEQ` unacknowledged frames

Notes

Go-back-n (2/2)

```
1  while (true) {
2      wait_for_event(&event);
3      switch(event) {
4          case network_layer_ready:
5              from_network_layer(&buffer[next_frame_to_send]);
6              nbuffered = nbuffered + 1;
7              send_data(next_frame_to_send, frame_expected, buffer);
8              inc(next_frame_to_send);
9              break;
10         case frame_arrival:
11             from_physical_layer(&r);
12             if (r.seq == frame_expected) {
13                 to_network_layer(&r.info);
14                 inc(frame_expected);
15             }
16             while (between(ack_expected, r.ack, next_frame_to_send)) {
17                 nbuffered = nbuffered - 1;
18                 stop_timer(ack_expected);
19                 inc(ack_expected);
20             }
21             break;
22         case cksum_err: break; /* just ignore bad frames */
23         case timeout: /* trouble; retransmit outstanding frames */
24             next_frame_to_send = ack_expected;
25             for (i = 1; i <= nbuffered; i++) {
26                 send_data(next_frame_to_send, frame_expected, buffer);
27                 inc(next_frame_to_send);
28             }
29             if (nbuffered < MAX_SEQ)
30                 enable_network_layer();
31             else
32                 disable_network_layer();
33     }
34 }
```

Notes

Go-back-n

Comments

- An ack for frame `n`, automatically acknowledges also frames `n-1`, `n-2`, etc.
- Note that a maximum of `MAX_SEQ` frames and not `MAX_SEQ + 1` frames may be outstanding
 - `packet buffer[MAX_SEQ+1]`
 - `if nbuffered < MAX_SEQ`
- The reason is the following:
 - 1 The sender sends frames 0 through 7
 - 2 An ack for frame 7 comes back to the sender
 - 3 The sender sends other eight frames with sequence number 0 through 7
 - 4 Now another piggybacked ack for frame 7 comes in
- **Problem:** Did all eight frames belonging to the second batch arrive correctly or did all get lost ?
 - the sender has no way to detect it
- **Solution** Restrict the maximum number of outstanding frames to `MAX_SEQ`

Notes

Selective Repeat

- Go-back-n works well if errors are rare, but if the line is poor, it wastes a lot of bandwidth
- An alternative strategy is to allow receiver to accept and buffer the frames following a damaged or lost one
 - both sender and receiver maintain a window larger than 1
 - whenever a frame arrives, its sequence number is checked by the function `between` to see if it falls within the window
 - however, it must be kept within the data link layer and not passed to the network layer until all lower numbered frames are received

Notes

Selective Repeat (1/3)

```
1 static boolean between(seq_nr a, seq_nr b, seq_nr c) {...}
2 static void send_frame(frame_kind fk, seq_nr frame_nr,
3   seq_nr frame_expected, packet buffer[]){
4   frame s; s.kind = fk;
5   if (fk == data) s.info = buffer[frame_nr % NR_BUFS];
6   s.seq = frame_nr;
7   s.ack = (frame_expected + MAX_SEQ) % (MAX_SEQ + 1);
8   if (fk == nak) no_nak = false;
9   to_physical_layer(s);
10  if (fk == data) start_timer(frame_nr % NR_BUFS);
11  stop_ack_timer();
12  {
13    void protocol6(void){
14      seq_nr ack_expected, next_frame_to_send, frame_expected;
15      seq_nr nbuffered, too_far; event_type event; int i; frame r;
16      packet out_buf[NR_BUFS], in_buf[NR_BUFS];
17      boolean arrived[NR_BUFS];
18
19      enable_network_layer();
20      ack_expected = 0; next_frame_to_send = 0; frame_expected = 0;
21      too_far = NR_BUFS; nbuffered = 0;
22      for (i = 0; i < NR_BUFS; i++) arrived[i] = false;
23      while (true) {
24        wait_for_event(sevent);
25        switch(event) {
26          case network_layer_ready:
27            nbuffered = nbuffered + 1;
28            from_network_layer(sout_buf[next_frame_to_send % NR_BUFS]);
29            send_frame(data, next_frame_to_send, frame_expected, out_buf);
30            inc(next_frame_to_send);
31            break;
```

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols

45 / 55

Notes

Example: Selective Repeat (2/3)

```
1 case frame_arrival:
2   from_physical_layer(r);
3   if (r.kind == data) {
4     if ((r.seq != frame_expected) && no_nak)
5       send_frame(nak, 0, frame_expected, out_buf);
6     else start_ack_timer();
7     if (between(frame_expected, r.seq, too_far) &&
8       if(arrived[r.seq % NR_BUFS] == false)) {
9       arrived[r.seq % NR_BUFS] = true;
10      in_buf[r.seq % NR_BUFS] = r.info;
11      while (arrived[frame_expected % NR_BUFS]) {
12        to_network_layer(sin_buf[frame_expected % NR_BUFS]);
13        no_nak = true;
14        arrived[frame_expected % NR_BUFS] = false;
15        inc(frame_expected);
16        inc(too_far);
17        start_ack_timer(); }
18    }
19    if ((r.kind == nak) &&
20      if(between(ack_expected, (r.ack+1) % (MAX_SEQ+1),
21        next_frame_to_send))
22      send_frame(data, (r.ack+1) % (MAX_SEQ + 1),
23        frame_expected, out_buf);
24
25    while (between(ack_expected, r.ack, next_frame_to_send)) {
26      nbuffered = nbuffered + 1;
27      stop_timer(ack_expected % NR_BUFS);
28      inc(ack_expected);
29    }
30    break;
```

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols

46 / 55

Notes

Example: Selective Repeat (3/3)

```
1 case cksun_err:
2   if (no_nak) send_frame(nak, 0, frame_expected, out_buf);
3   break;
4
5 case timeout:
6   send_frame(data, oldest_frame, frame_expected, out_buf);
7   break;
8
9 case ack_timeout:
10  send_frame(ack, 0, frame_expected, out_buf);
11  }
12  if (nbuffered < NR_BUFS) enable_network_layer();
13  else disable_network_layer();
14  }
15 }
```

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols


47 / 55

Notes

Selective Repeat

Comments

- Frames need not be received in order,
 - we may have an undamaged frame N , while still waiting for an undamaged version of $N - 1$.
- If the receiver delivers all frames in its window just after sending an ack for the entire window, we may have a serious problem:



- Solution:** we must avoid overlapping send and receive windows
⇒ the highest sequence number must be at least twice the window size.

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols

48 / 55

Notes

Data Link Protocols

Demo

http://www.csi.uottawa.ca/~elsaddik/abedweb/applets/Applets/Sliding_Window/sliding-window/index.html

Paolo Costa

03 - Data Link Layer

Sliding Window Protocols 49 / 55

Data Link Layer Protocols

Now let's take a look how point-to-point connections are supported in, for example, the Internet.

Recall:

- The data link layer is responsible for transmitting **frames** from sender to receiver.
- It can use only the physical layer, which supports only transmission of a bit at a time.
- The DLL has to take into account that transmission errors may occur
 - ⇒ **error control** (ACKs, NACKs, checksums, etc.)
- The DLL has to take into account that sender and receiver may operate at different speeds
 - ⇒ **flow control** (windows, frame numbers, etc.)

Paolo Costa

03 - Data Link Layer

Examples 50 / 55

High-Level Data Link Control

- **HDLC**: A pretty old, but widely used protocol for point-to-point connections. Is **bit-oriented**.

Bits	8	8	8	> 0	16	8
	0 1 1 1 1 1 0	Address	Control	Data	Checksum	0 1 1 1 1 1 0

- HDLC uses a sliding window protocol with 3-bit sequencing
- The **control** field is used to distinguish different kinds of frames
 - contains sequence numbers, acks, nacks, etc.

Question

- What do we need the **address** field for? **Useful on line with multiple terminals**


Paolo Costa

03 - Data Link Layer

Examples 51 / 55

Internet Point-to-Point Connections

- This is what may happen when you have an Internet connection to a provider:



- We'd like to use the Internet protocol stack at our home (e.g., ADSL)
 - the bottom line is that we'll have to transfer IP (network) packets across our telephone line
- **Issues:**
 - how can we embed IP packets into frames ?
 - how can we unpacked at the other end ?

Paolo Costa

03 - Data Link Layer

Examples 52 / 55

Notes

Notes


Notes

Notes

PPP: Point-to-Point Protocol

- **PPP** is the data link protocol for point-to-point connections with respect to the Internet (**home-ISP** and **router-to-router**):
 - **Proper framing**, i.e. the start and end of a frame can be unambiguously detected.
 - A separate protocol for **controlling the line** (setup, testing, negotiating options, and tear-down) (**LCP**)
 - Supports **many different network layer protocols**, not just IP.
 - No need for **fixed network addresses**.

The default frame:



Address (no data link addresses) **Control** no seq. number (no reliable transmission)
Protocol LCP / NCP / IP / ... **Payload** (up to a negotiated maximum)
Checksum (usually 2 bytes)

Paolo Costa 03 - Data Link Layer Examples 53 / 55

Notes

PPP

Connection Setup

Suppose you want to set up a true Internet connection to your provider.


- 1 **Set up a physical connection** through your modem
- 2 Your PC starts sending a number of **Link Control Packets** (LCP) to **negotiate the kind of PPP connection** you want.
 - the maximum payload size in data frames
 - do authentication (e.g. ask for a password)
 - monitor the quality of the link (e.g. how many frames didn't come through).
 - compress headers (useful for slow links between fast computers)
- 3 Then, we **negotiate network layer stuff**, like getting an IP address that the provider's router can use to forward packets to you.

Paolo Costa 03 - Data Link Layer Examples 54 / 55

Notes

PPP

Example



Question

- If an IP address is dynamically assigned, who does the assignment ? **The provider**

Question

- If an IP address is dynamically assigned, Can someone else ever send you data (they don't know your address, do they?) **We need to contact them first (our address is included in the request).**

Paolo Costa 03 - Data Link Layer Examples 55 / 55

Notes

Notes
